

<https://ukrainer.net/en/>

[https://www.instagram.com/ukrainer\\_net/](https://www.instagram.com/ukrainer_net/)

# DSCI 510

# Principles of Programming

# for Data Science

**Jose-Luis Ambite**

Research Team Leader, Information Sciences Institute  
Associate Research Professor, Department of Computer Science

University of Southern California

# **Databases and SQL**

## **Chapter 15 in textbook**

# "DB" Classes

- DSCI-551 Foundations of Data Management
- DSCI-550 Data Science at Scale
- DSCI-558/CSCI-563 Building Knowledge Graphs
- CSCI-585 Database Systems
- CSCI-548 Information Integration on the Web

# Why Databases?

When you have lots of data (i. e., too much to store in main memory), you need to:

- 1) organize your data
- 2) be able to find what you need quickly

# Database

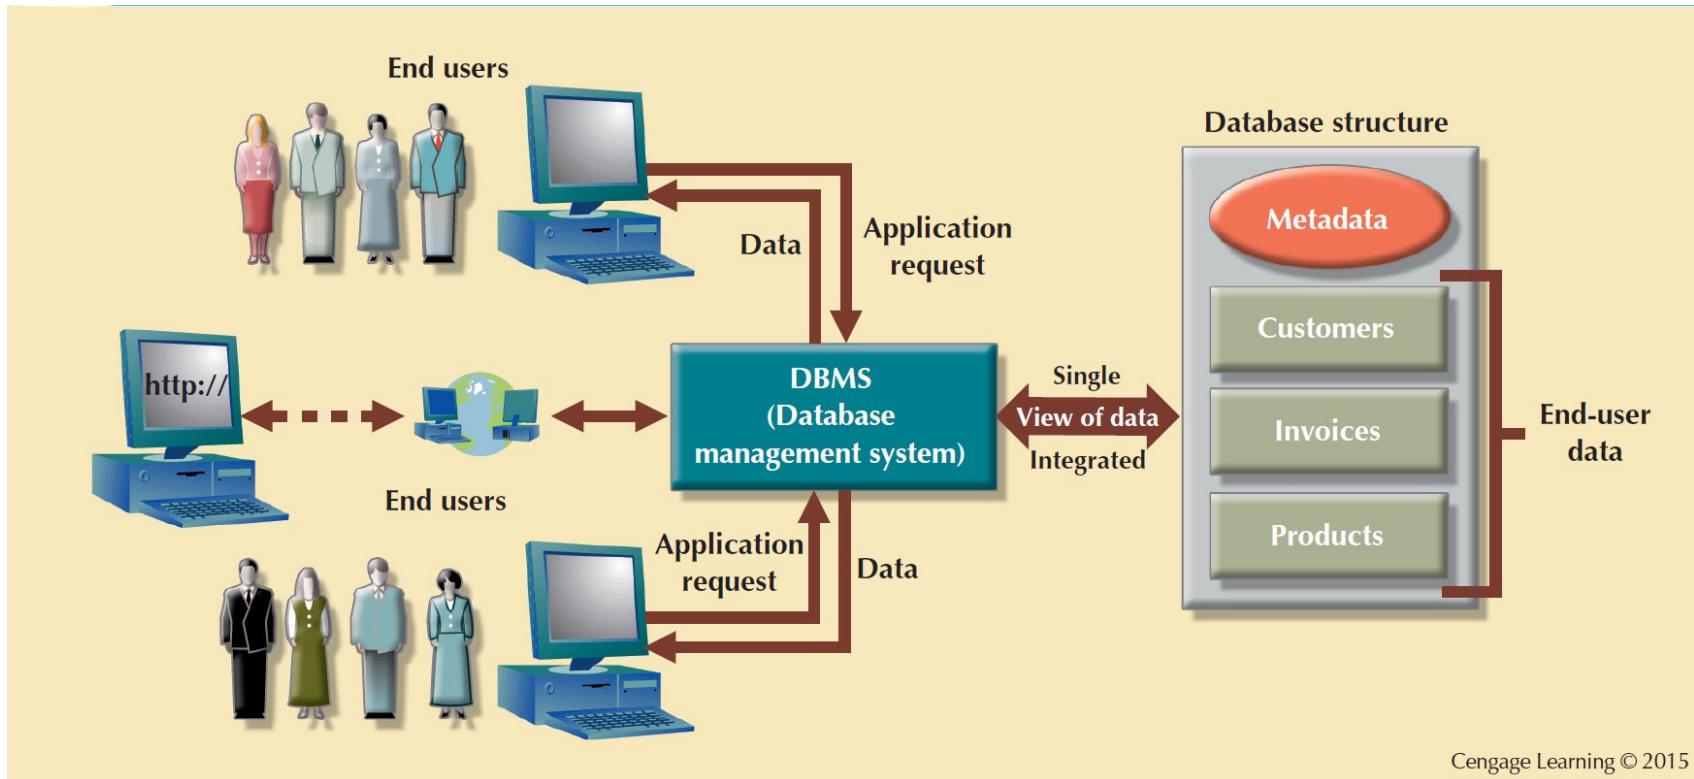
- Logically coherent collection of data with inherent meaning
- Represents some aspect of the real world
- Miniworld or universe of discourse
  - Logic: FOL Model
- Built for a specific purpose
- Maintains metadata (data about data): describes data characteristics and relationships

# Database Management System (DBMS)

Collection of programs that:

- Manage the database structure
- Evaluate queries efficiently against the data
- Control access to data stored in the database
- Manage concurrent access/modifications by multiple users

# DBMS Manages Interaction between End User and Database



# DBMS Functions

## Data dictionary management

- **Data dictionary:** Stores definitions of the data elements and their relationships

## Data storage management

- **Performance tuning:** Ensures efficient performance of the database in terms of storage and access speed

## Data transformation and presentation

- Transforms entered data to conform to required data structures

## Security management

- Enforces user security and data privacy

# DBMS Functions

## Multiuser access control

- Sophisticated algorithms ensure that multiple users can access the database concurrently without compromising its integrity

## Backup and recovery management

- Enables recovery of the database after a failure

## Data integrity management

- Minimizes redundancy and maximizes consistency

# DBMS Functions

Database access languages and application programming interfaces

- **Query language:** Lets the user specify what must be done without having to specify how
- **Structured Query Language (SQL):** De facto query language and data access standard supported by the majority of DBMS vendors

Database communication interfaces

- Accept end-user requests via multiple, different network environments

# Relational Databases

Relational databases model data by storing rows and columns in tables. The power of the relational database lies in its ability to efficiently retrieve data from those tables and in particular where there are multiple tables and the relationships between those tables involved in the query.

[http://en.wikipedia.org/wiki/Relational\\_database](http://en.wikipedia.org/wiki/Relational_database)

# Terminology

- Database - contains many tables
- Relation (or table) - contains tuples and attributes
- Tuple (or row) - a set of fields that generally represents an “object” (entity) like a person or a music track
- Attribute (also column or field) - one of possibly many elements of data corresponding to the object represented by the row

# Relation, Database (Discrete Math)

Given  $k$  sets (domains)  $D_1, D_2, \dots, D_k$ ,

the **cartesian product** of these sets  $D_1 \times D_2 \times \dots \times D_k$

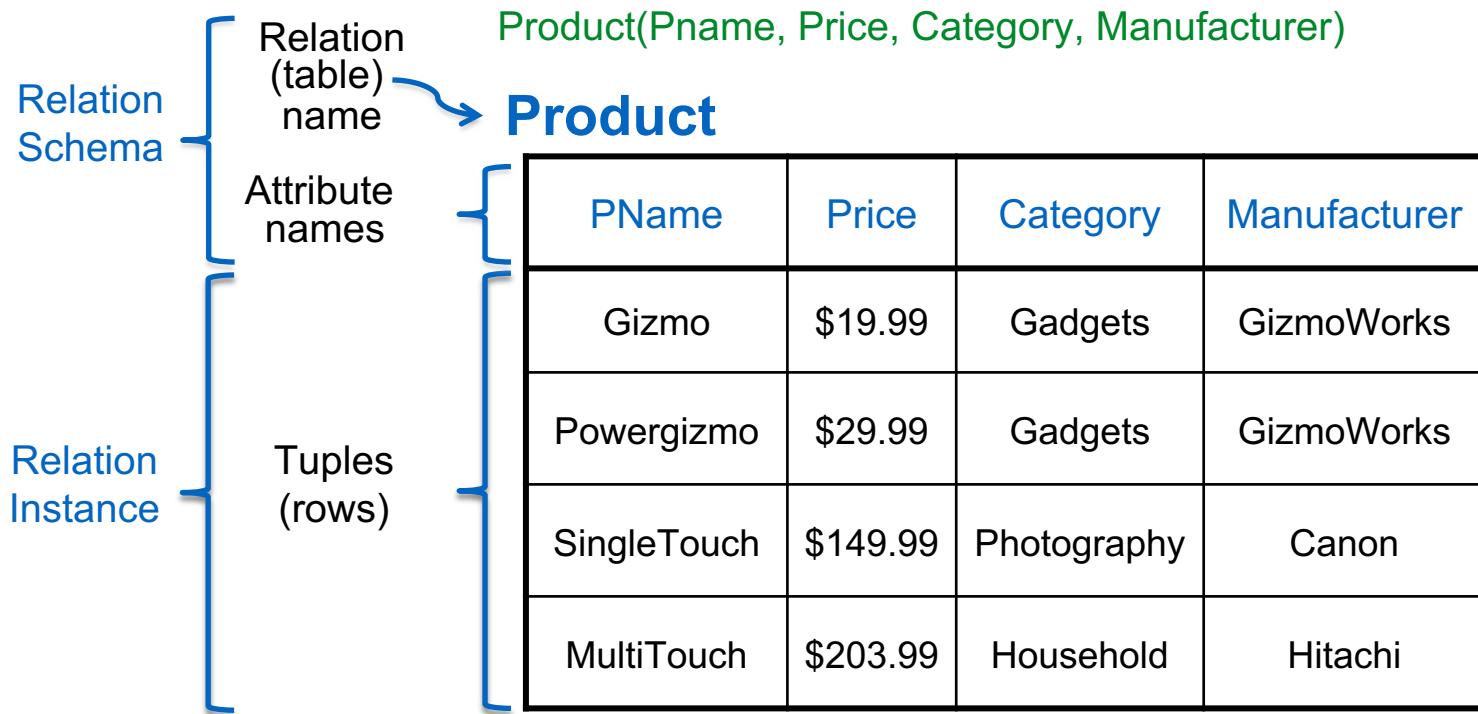
is the set of all  $k$ -tuples  $(d_1, d_2, \dots, d_k)$  such that  $d_i \in D_i$  for  $1 \leq i \leq k$

A  $k$ -ary **relation  $R$**  is a subset of the cartesian product of  $k$  sets, i.e.,

$$R \subseteq D_1 \times D_2 \times \dots \times D_k$$

A **database** is a set of relations

# Relational schema, instance



**STUDENT**

Name	Student_number	Class	Major
Smith	17	1	CS
Brown	8	2	CS

**COURSE**

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

**SECTION**

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	07	King
92	CS1310	Fall	07	Anderson
102	CS3320	Spring	08	Knuth
112	MATH2410	Fall	08	Chang
119	CS1310	Fall	08	Anderson
135	CS3380	Fall	08	Stone

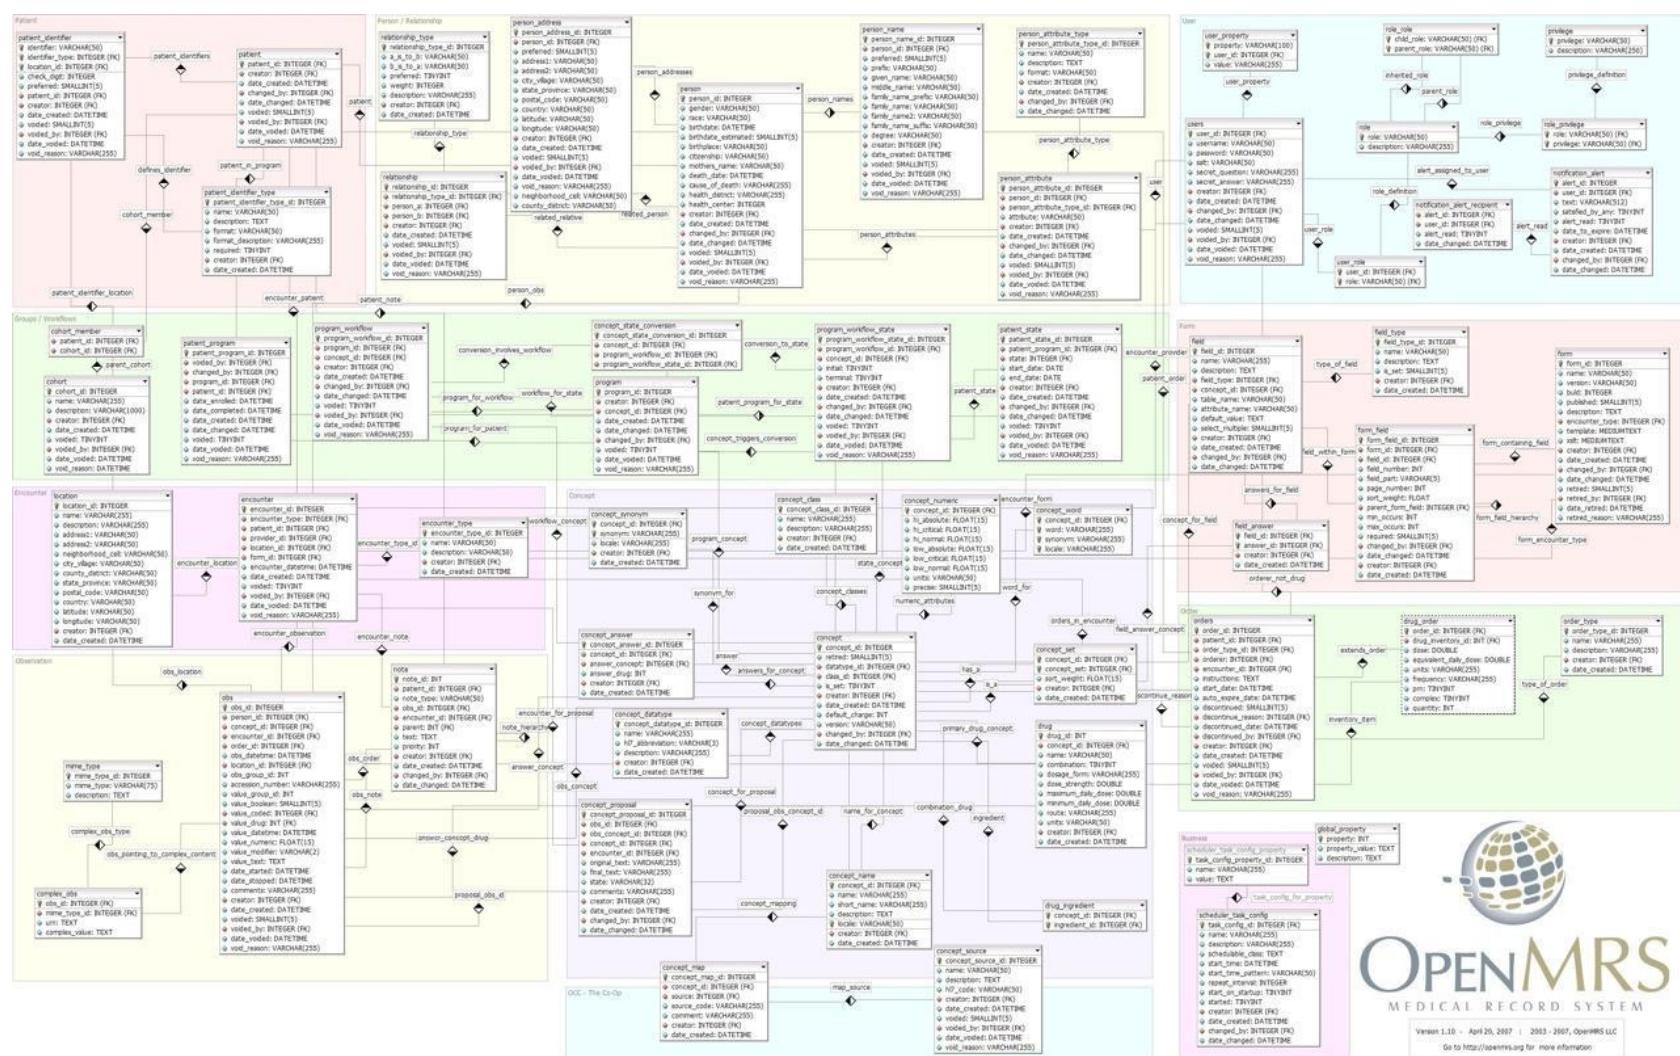
**Figure 1.2**  
A database that stores student and course information.

**GRADE REPORT**

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

**PREREQUISITE**

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310



# OPENMRS

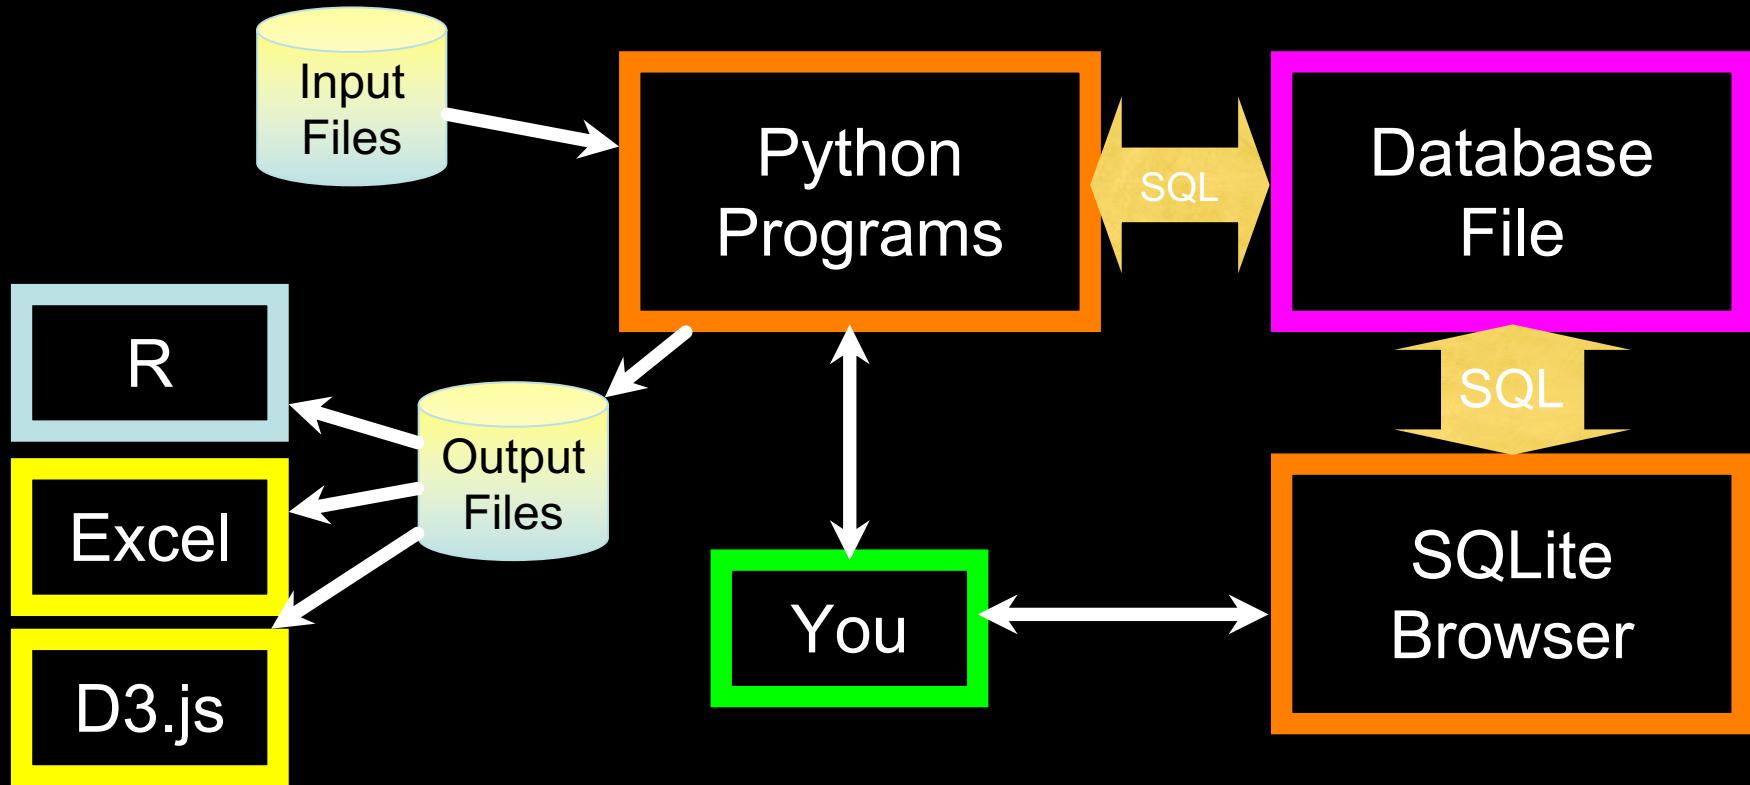
MEDICAL RECORD SYSTEM

# SQL

Structured Query Language is the language we use to issue commands to the database

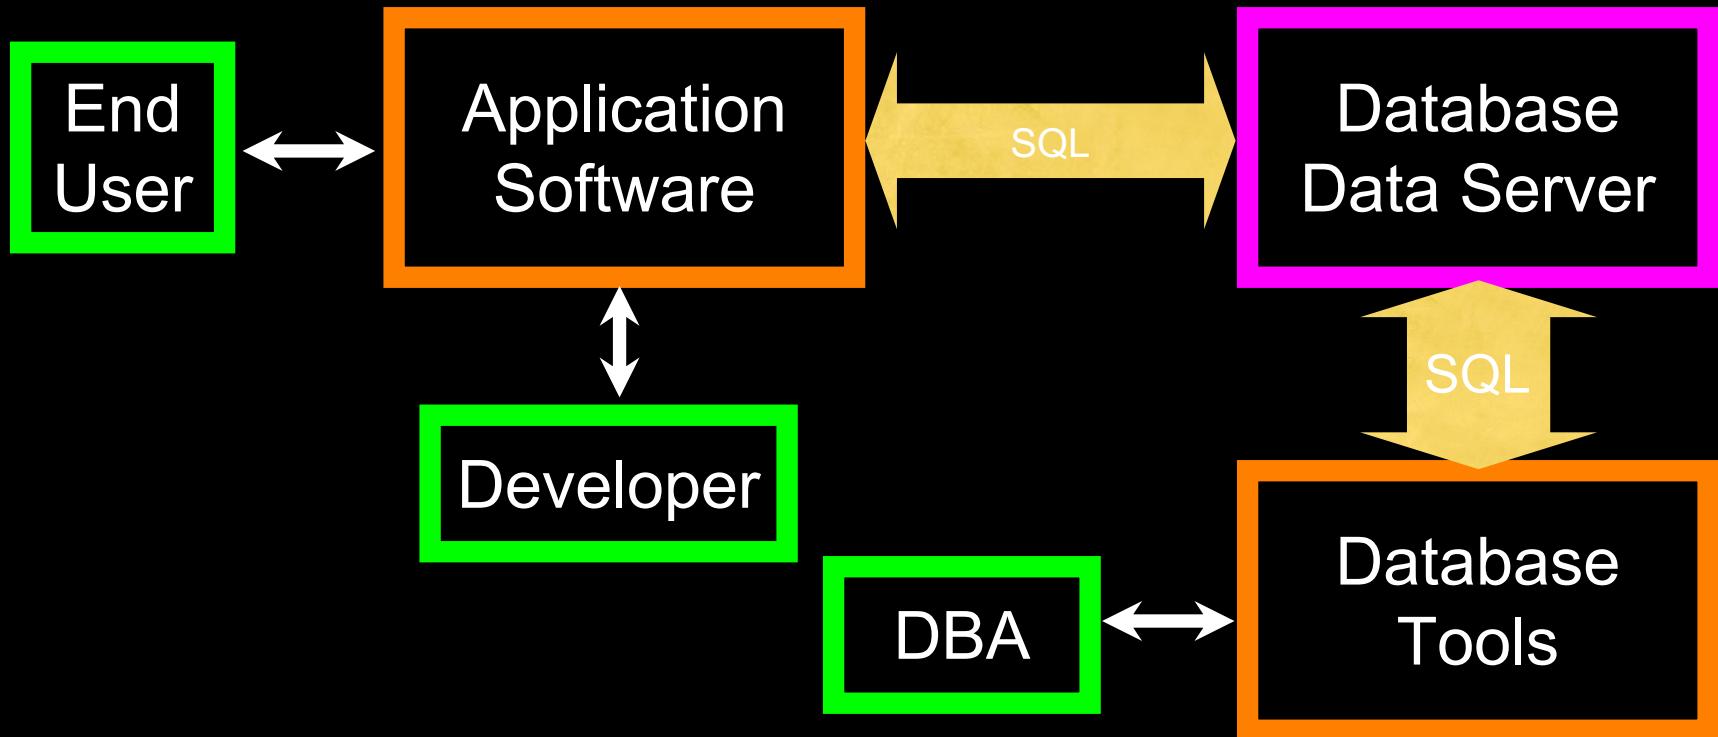
- Create data (a.k.a. Insert)
- Retrieve data
- Update data
- Delete data

<http://en.wikipedia.org/wiki/SQL>



# Major Web Applications w/ Databases

- **Application Developer** - Builds the logic for the application, the look and feel of the application - monitors the application for problems
- **Database Administrator** - Monitors and adjusts the database as the program runs in production
- Often both people participate in the building of the “Data model”



DBA = Database Administrator

# Database Administrator

A database administrator (DBA) is a person responsible for the design, implementation, maintenance, and repair of an organization's database. The role includes the development and design of database strategies, monitoring and improving database performance and capacity, and planning for future expansion requirements. They may also plan, coordinate, and implement security measures to safeguard the database.

[http://en.wikipedia.org/wiki/Database\\_administrator](http://en.wikipedia.org/wiki/Database_administrator)

# Common Database Systems

- Three major Database Management Systems in wide use
  - Oracle - Large, commercial, enterprise-scale, very very tweakable
  - MySQL - Simpler but very fast and scalable - commercial open source
  - SQLServer - Very nice - from Microsoft (also Access)
- Many other smaller projects, free and open source
  - HSQL, SQLite, PostgreSQL (aka Postgres), ...

We will be  
using this

But for anything more serious  
you should probably use  
PostgreSQL or MySQL

# SQLite

- SQLite is a very popular database - it is free and fast and small
- SQLite Browser allows us to directly manipulate SQLite files
  - <http://sqlitebrowser.org/>
- SQLite is embedded in Python and a number of other languages

# SQLite Is in Lots of Software...

symbian

python™



skype™



Microsoft®

McAfee®



A®  
Adobe



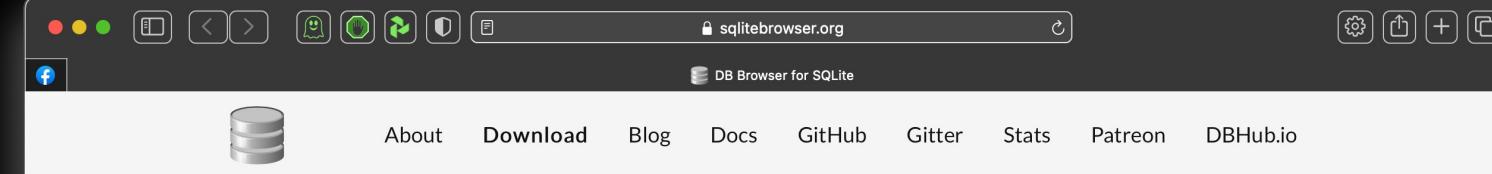
php

Google™

TOSHIBA

Sun  
microsystems

<http://www.sqlite.org/famous.html>



# DB Browser for SQLite

*The Official home of the DB Browser for SQLite*

## Screenshot

A screenshot of the SQLite Database Browser application. The title bar says "SQLite Database Browser - /Users/jc/tmp/example.db". The menu bar includes "File", "Edit", "View", "Tools", "Help", and "About". Below the menu is a toolbar with buttons for "New Database", "Open Database", "Write Changes", and "Revert Changes". There are four tabs: "Database Structure", "Browse Data" (which is selected), "Edit Pragmas", and "Execute SQL". Under "Browse Data", there is a "Table:" dropdown set to "total\_members" and buttons for "New Record" and "Delete Record". The main area shows a table with three columns: "list", "month", and "members". The data is as follows:

	list	month	members
1	gluster-board	2013-09-05	99999
2	gluster-users	2013-09-05	99999

< 1 - 2 of 12 >

Go to:

1

SQL Log

<http://sqlitebrowser.org/>

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

Create Table Create Index Modify Table Delete Table

Name	Type	Schema
Tables (0)		
Indices (0)		
Views (0)		
Triggers (0)		

<http://sqlitebrowser.org/>

Edit Database Cell

Mode: Text Import Export Set as NULL

Type of data currently in cell: NULL  
0 byte(s)

DB Schema

Name	Type	Schema
Tables (0)		
Indices (0)		
Views (0)		
Triggers (0)		

SQL Log Plot DB Schema Remote

# Lets Make a Database

<https://www.py4e.com/lectures3/Pythonlearn-15-Database-Handout.txt>

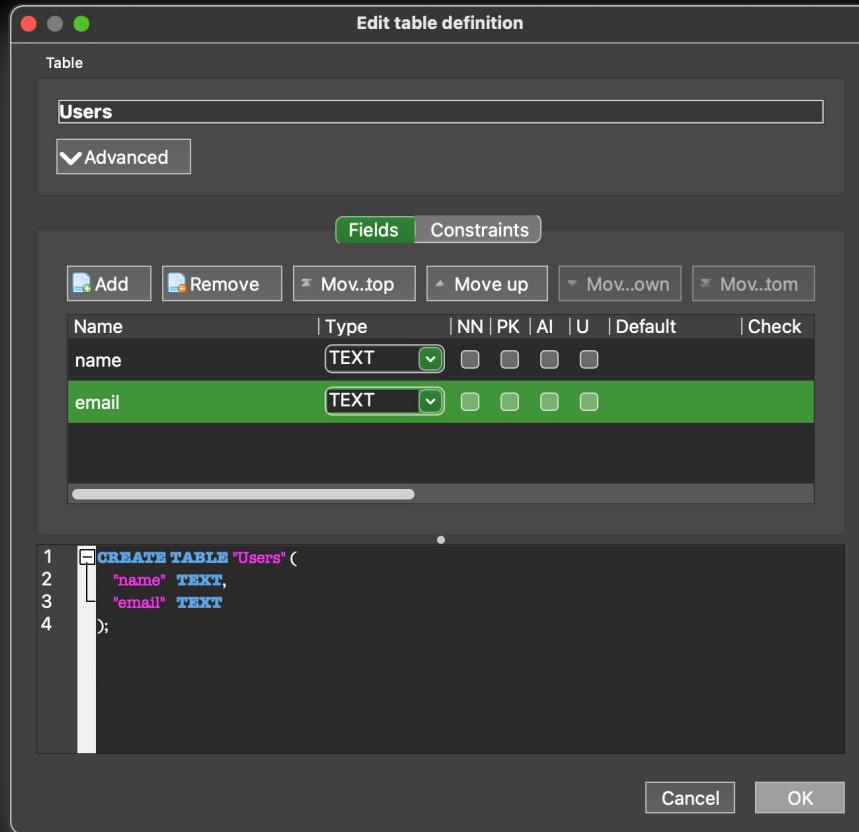
Go to week12-Database-Handout.txt  
Show in SQLite browser

# Please Note!

The following slides describe what we can do with the DB Browser application. It's not practical to constantly shift back and forth between my slides and the actual application.

When you get home, download and open the DB Browser application and go through all the steps I did to make sure you understand **what** I did and **how to use** the database application.

# Start Simple - A Single Table



```
CREATE TABLE Users (
    name TEXT,
    email TEXT
)
```

File New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Create Table Create Index Print

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )
name	TEXT	"name" TEXT
email	TEXT	"email" TEXT
Indices (0)		
Views (0)		
Triggers (0)		

Mode: Text

1

Type of data currently in cell

Size of data currently in table

DB Schema

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )
name	TEXT	"name" TEXT
email	TEXT	"email" TEXT
Indices (0)		
Views (0)		
Triggers (0)		

SQL Log Plot DB Schema Remote

UTF-8

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Table: Users Filter in ...

	name	email
1	Chuck	csev@umich.edu
2	Colleen	cvl@umich.edu
3	Ted	ted@umich.edu
4	Sally	a1 @umich.edu

Our table with four rows

Mode: Text

1 a1 @umich.edu

Type of data currently in cell: Text / Numeric  
12 character(s) Apply

DB Schema

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT,"email" TEXT )
Indices (0)		
Views (0)		
Triggers (0)		

Go to: 1

SQL Log Plot DB Schema Remote

UTF-8

# SQL

Structured Query Language is the language we use to issue commands to the database

- Create data (a.k.a Insert)
- Retrieve data
- Update data
- Delete data

<http://en.wikipedia.org/wiki/SQL>

# SQL: Insert

The Insert statement inserts a row into a table

```
INSERT INTO Users (name, email) VALUES ('Kristin', 'kf@umich.edu')
```

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Mode: Text

1 `INSERT INTO Users (name, email) VALUES ('Kristin', 'kf@umich.edu')`

Type of data currently in cell: NULL  
0 byte(s) Apply

DB Schema

Name	Type	Schema
Tables (1)		
Users	CREATE TABLE "Users" ( "name" TEXT,"email" TEXT )	
Indices (0)		
Views (0)		
Triggers (0)		

Execution finished without errors.  
Result: query executed successfully. Took 0ms, 1 rows affected  
At line 1:  
`INSERT INTO Users (name, email) VALUES ('Kristin', 'kf@umich.edu')`

SQL Log Plot DB Schema Remote

UTF-8

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

DB Browser for SQLite - /Users/arens/Desktop/NewDB.db

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Table: Users

	name	email
1	Ted	ted@umich.edu
2	Sally	a1 @umich.edu
3	Kristin	kf@umich.edu
4	Colleen	cvl@umich.edu
5	Chuck	csev@umich.edu

1 INSERT ID

Execution finished  
Result: query executed  
At line 1:  
INSERT INTO Us

Go to: 1

Filter in ...

Mode: Text

1 3

Type of data currently in cell: Text / Numeric  
1 character(s) Apply

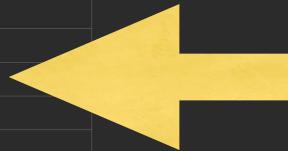
DB Schema

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )
Indices (0)		
Views (0)		
Triggers (0)		

SQL Log Plot DB Schema Remote

1 - 5 of 5 Go to: 1

UTF-8



# SQL: Delete

Deletes a row in a table based on selection criteria

```
DELETE FROM Users WHERE email='ted@umich.edu'
```

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Mode: Text

1 `DELETE FROM Users WHERE email='ted@umich.edu'`

Execution finished without errors.  
Result: query executed successfully. Took 0ms, 1 rows affected  
At line 1:  
`DELETE FROM Users WHERE email='ted@umich.edu'`

NULL

Type of data currently in cell: NULL  
0 byte(s)

DB Schema

Name	Type	Schema
Tables (1)		CREATE TABLE "Users" ( "name" TEXT,"email" TEXT )
> Users		
Indices (0)		
Views (0)		
Triggers (0)		

SQL Log Plot DB Schema Remote

UTF-8

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL DB Browser for SQLite - /Users/arenz/Desktop/NewDB.db Edit Database Cell

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

1 [DELETE] Database Structure Browse Data Edit Pragmas Execute SQL

Table: Users Filter Filter

	name	email
1	Sally	a1 @umich.edu
2	Kristin	kf@umich.edu
3	Colleen	cvl@umich.edu
4	Chuck	csev@umich.edu

Execution finish  
Result: query ex  
At line 1:  
DELETE FROM

1 - 4 of 4 Go to: 1

Mode: Text

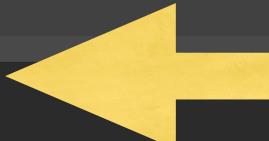
1 4

Type of data currently in cell: Text / Numeric  
1 character(s) Apply

DB Schema

Name	Type	Schema
Tables (1)		CREATE TABLE "Users" ( "name" TEXT,"email" TEXT )
Users		
Indices (0)		
Views (0)		
Triggers (0)		

SQL Log Plot DB Schema Remote



UTF-8

# SQL: Update

Allows the updating of a field with a where clause

```
UPDATE Users SET name='Charles'  
WHERE email='csev@umich.edu'
```

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Mode: Text

NULL

Type of data currently in cell: NULL  
0 byte(s)

DB Schema

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT,"email" TEXT )
Indices (0)		
Views (0)		
Triggers (0)		

Execution finished without errors.  
Result: query executed successfully. Took 0ms, 1 rows affected  
At line 1:  
UPDATE Users SET name='Charles' WHERE email='csev@umich.edu'

SQL Log Plot DB Schema Remote

UTF-8

New Database

DB Browser for SQLite - /Users/arenz/Desktop/NewDB.db

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Table: Users

name email

1 Sally a1 @umich.edu

2 Kristin kf@umich.edu

3 Colleen cvl@umich.edu

4 Charles csev@umich.edu

Filter Filter

Execution finished  
Result: query  
At line 1:  
UPDATE User

1 - 4 of 4 Go to: 1

Mode: Text

1 4

Type of data currently in cell: Text / Numeric  
1 character(s)

Apply

DB Schema

Name Type Schema

Tables (1)  
Users  
Indices (0)  
Views (0)  
Triggers (0)

CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )

SQL Log Plot DB Schema Remote

UTF-8

# Retrieving Records: Select

The select statement retrieves a group of records – you can either retrieve all the records or a subset of the records with a WHERE clause

```
SELECT * FROM Users
```

\* means retrieve all the columns

```
SELECT email FROM Users
```

or just list the columns you want to get

```
SELECT * FROM Users  
WHERE email='csev@umich.edu'
```

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Mode: Text

NULL

Type of data currently in cell: NULL

0 byte(s)

DB Schema

Name	Type	Schema
Tables (1)		
Users	CREATE TABLE "Users" ( "name" TEXT,"email" TEXT )	
Indices (0)		
Views (0)		
Triggers (0)		

Execution finished without errors.  
Result: 4 rows returned in 7ms  
At line 1:  
SELECT \* FROM Users

SQL Log Plot DB Schema Remote

UTF-8

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

S... S...

1 `SELECT * FROM Users WHERE email='csev@umich.edu'`

name	email
Charles	csev@umich.edu

Execution finished without errors.  
Result: 1 rows returned in 6ms  
At line 1:  
`SELECT * FROM Users WHERE email='csev@umich.edu'`

Edit Database Cell

Mode: Text

NULL

Type of data currently in cell: NULL

0 byte(s)

Apply

DB Schema

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )
Indices (0)		
Views (0)		
Triggers (0)		

SQL Log Plot DB Schema Remote

UTF-8

# Sorting with ORDER BY

You can add an ORDER BY clause to SELECT statements to get the results sorted in ascending or descending order

```
SELECT * FROM Users ORDER BY email
```

Descending  
order

```
SELECT * FROM Users ORDER BY name DESC
```

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Mode: Text

NULL

Type of data currently in cell: NULL

0 byte(s)

DB Schema

Name	Type	Schema
Tables (1)		
Users		CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )
Indices (0)		
Views (0)		
Triggers (0)		

Execution finished without errors.  
Result: 4 rows returned in 6ms  
At line 1:  
SELECT \* FROM Users ORDER BY email

SQL Log Plot DB Schema Remote

UTF-8

New Database Open Database Write Changes Revert Changes Open Project Save Project Attach Database Close Database

Database Structure Browse Data Edit Pragmas Execute SQL

Mode: Text

`SELECT * FROM Users ORDER BY name DESC`

NULL

Type of data currently in cell: NULL

0 byte(s)

DB Schema

Name	Type	Schema
Tables (1)		
Users	CREATE TABLE "Users" ( "name" TEXT, "email" TEXT )	
Indices (0)		
Views (0)		
Triggers (0)		

Execution finished without errors.  
Result: 4 rows returned in 6ms  
At line 1:  
`SELECT * FROM Users ORDER BY name DESC`

Descending

SQL Log Plot DB Schema Remote

UTF-8

# SQL Summary

```
INSERT INTO Users (name, email) VALUES ('Kristin', 'kf@umich.edu')

DELETE FROM Users WHERE email='ted@umich.edu'

UPDATE Users SET name="Charles" WHERE email='csev@umich.edu'

SELECT * FROM Users

SELECT name, email FROM Users

SELECT * FROM Users WHERE email='csev@umich.edu'

SELECT * FROM Users ORDER BY email
```

# Creating a Database in Python

```
import sqlite3

# create a connection to a database (file)
# If database file does not exist, sqlite3 creates it for you
conn = sqlite3.connect('emaildb.sqlite') # emaildb.sqlite is a filename

# A database cursor is a mechanism that enables traversal over the records in a database.
# It allows retrieval, addition and removal of database records.
# It's an iterator.
cur = conn.cursor()

cur.execute('DROP TABLE IF EXISTS Counts')
cur.execute('CREATE TABLE Counts (email TEXT, count INTEGER)')

with open('mbox-short.txt') as f:
    for line in f:
        if not line.startswith('From: '): continue
        email = line.split()[1]
        cur.execute(f"SELECT count FROM Counts WHERE email = '{email}'")
        row = cur.fetchone()
        if row is None:
            cur.execute(f"INSERT INTO Counts (email, count) VALUES ('{email}', 1)")
        else:
            cur.execute(f"UPDATE Counts SET count = count + 1 WHERE email = '{email}'")
    conn.commit()

query = 'SELECT email, count FROM Counts ORDER BY count DESC LIMIT 10'

# cur.execute(query) returns an iterator over the query results
for row in cur.execute(query):
    print(f'{row[0]:32s} {row[1]:4d}')

cur.close() # we are done sending commands to the DB for now
conn.close() # we are completely done with the DB in this program
```

Counting emails by sender  
Go to week12-email-db.ipynb

# This is not too exciting (so far)

- Tables pretty much look like big fast programmable spreadsheets with rows, columns, and commands
- The power comes when we have more than one table and we can exploit the relationships between the tables

# Complex Data Models and Relationships

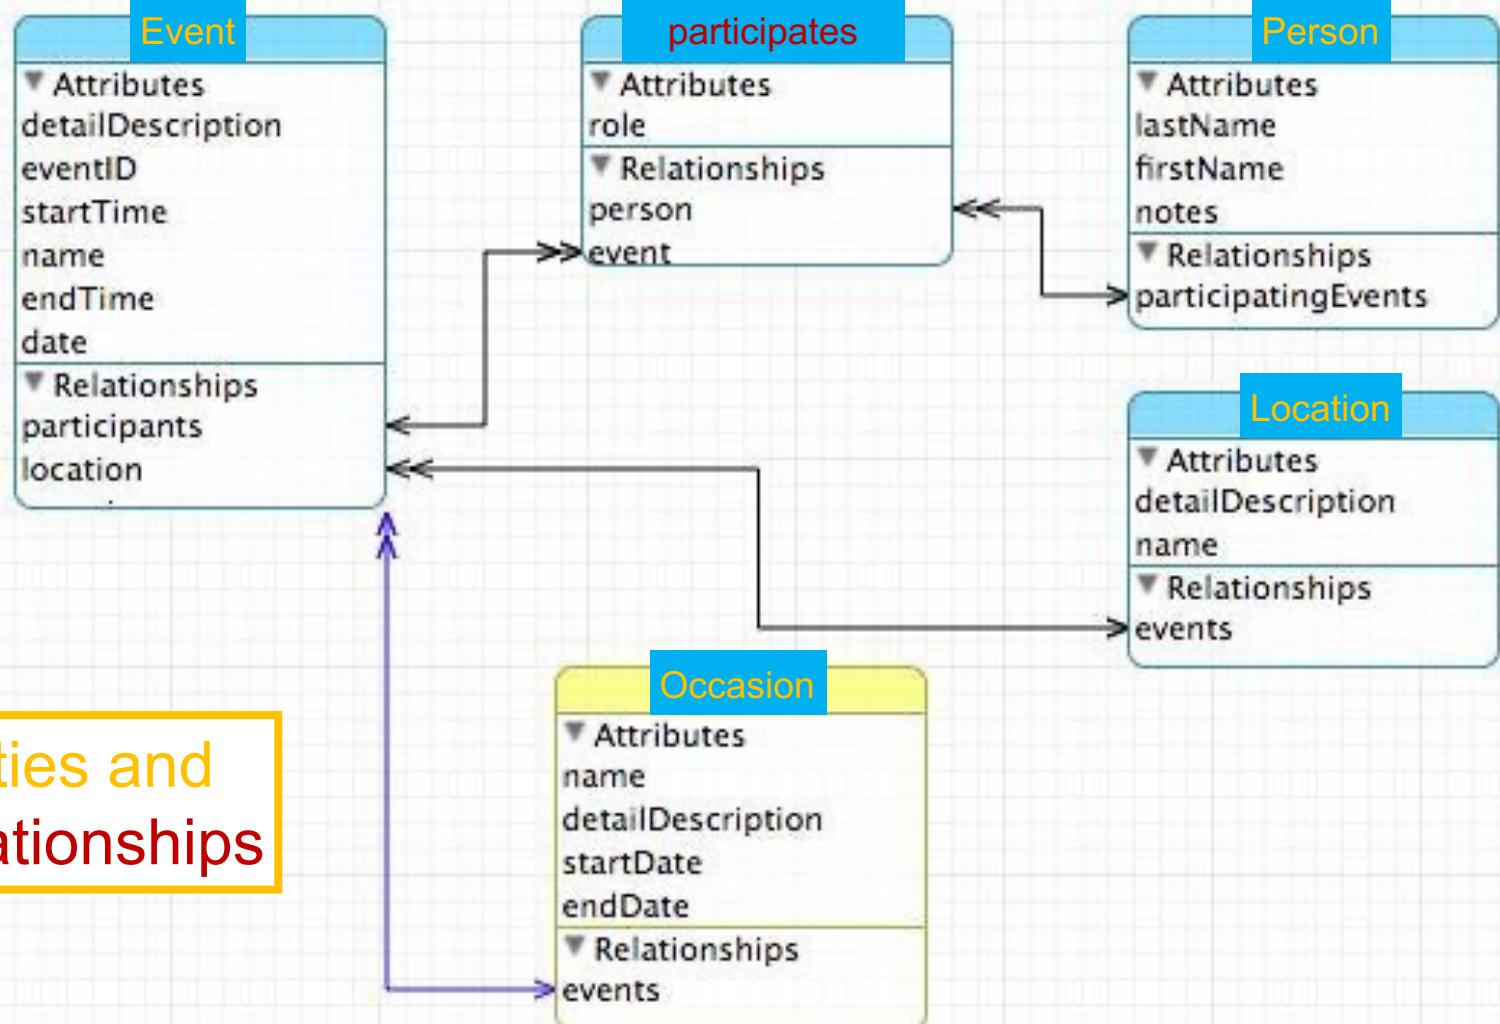
[http://en.wikipedia.org/wiki/Relational\\_model](http://en.wikipedia.org/wiki/Relational_model)

# Database Design

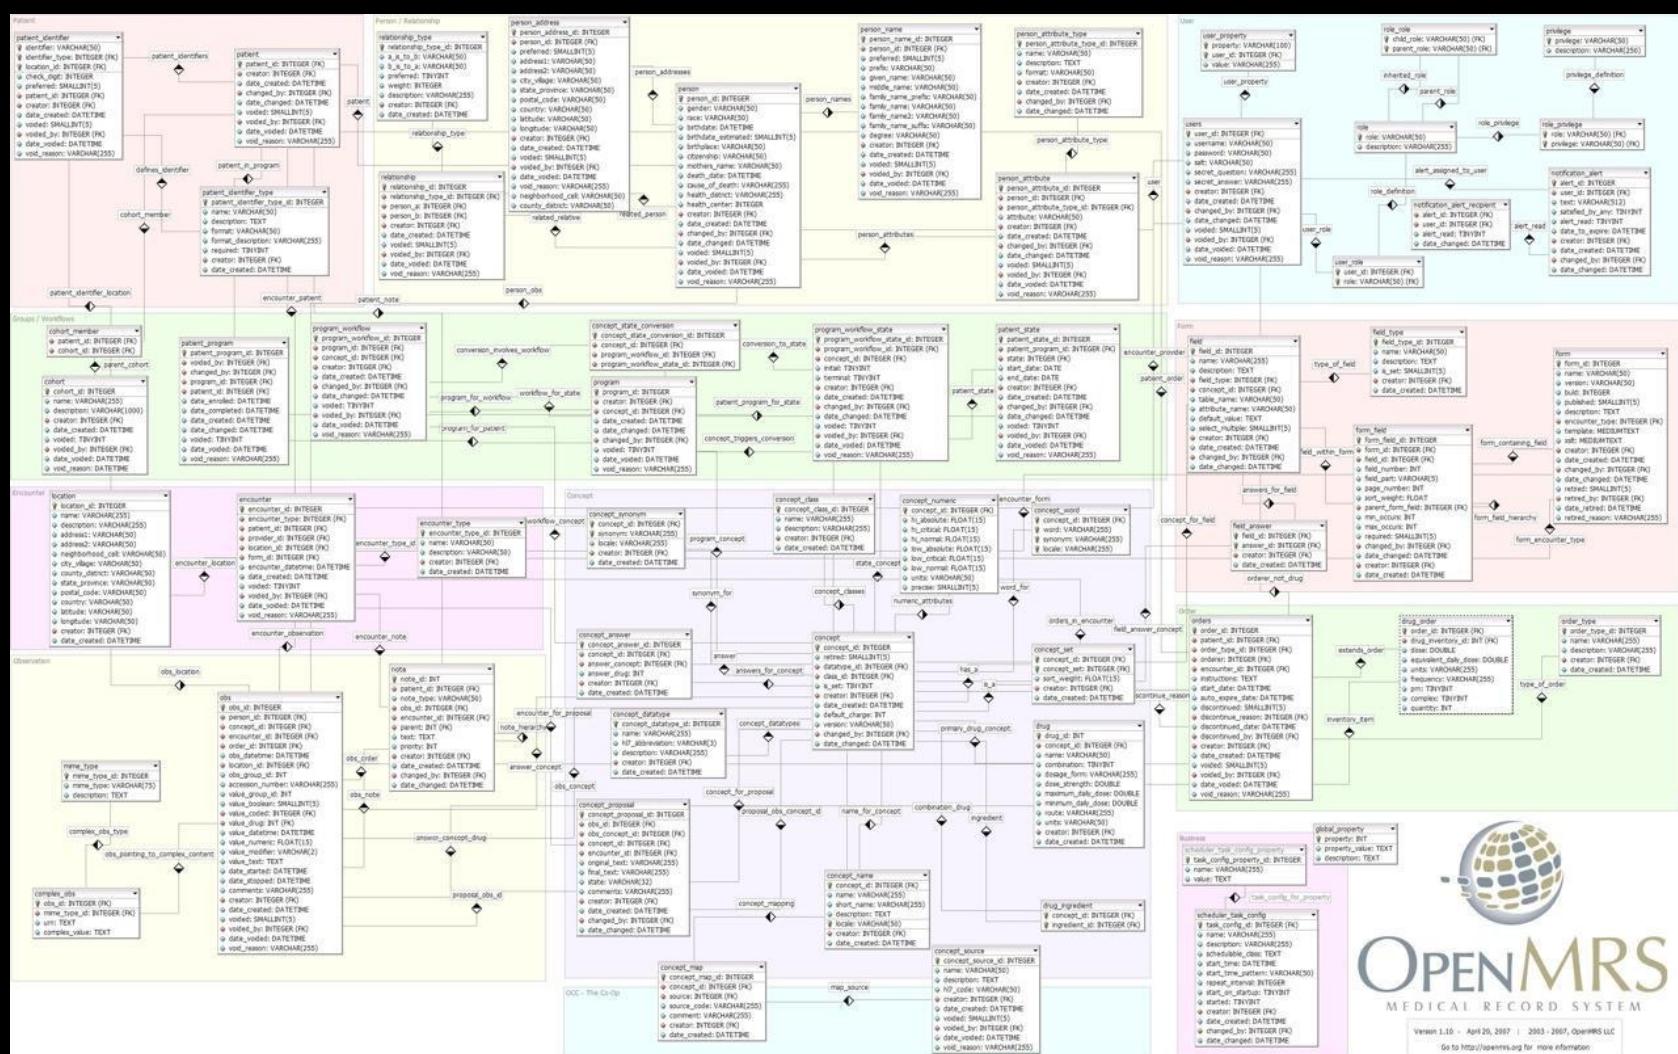
- Database design is an **art form** of its own with particular skills and experience
- Our goal is to avoid the really bad mistakes and design clean and easily understood databases
- Others may performance tune things later
- Database design starts with a picture...

# Building an (Entity-Relationship) Data Model

- Think of the entities (and their attributes) and the relationships among entities that exist in your domain/application
- Draw a picture of the entities and their relationships
- Usually entities are "nouns" and relationships "verbs"
- These entities and relationships turn into tables
- Basic Rule:
  - Don't put the same piece of data in twice - use a relationship instead
  - When there is one thing in the “real world” there should be one copy of that thing in the database



## Entities and Relationships



# OPENMRS

MEDICAL RECORD SYSTEM

Track	Len	Artist	Album	Genre	Rating	Count
<input checked="" type="checkbox"/> Helis Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tin Man	3:30	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Sister Golden Hair	3:22	America	Greatest Hits	Easy Listen...	★★★★★	24
<input checked="" type="checkbox"/> Track 01	4:22	Billy Price	Danger Zone	Blues/R&B	★★★★★	26
<input checked="" type="checkbox"/> Track 02	2:45	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
<input checked="" type="checkbox"/> Track 03	3:26	Billy Price	Danger Zone	Blues/R&B	★★★★★	22
<input checked="" type="checkbox"/> Track 04	4:17	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
<input checked="" type="checkbox"/> Track 05	3:50	Billy Price	Danger Zone	Blues/R&B	★★★★★	21
<input checked="" type="checkbox"/> War Pigs/Luke's Wall	7:58	Black Sabbath	Paranoid	Metal	★★★★★	25
<input checked="" type="checkbox"/> Paranoid	2:53	Black Sabbath	Paranoid	Metal	★★★★★	22
<input checked="" type="checkbox"/> Planet Caravan	4:35	Black Sabbath	Paranoid	Metal	★★★★★	25
<input checked="" type="checkbox"/> Iron Man	5:59	Black Sabbath	Paranoid	Metal	★★★★★	26
<input checked="" type="checkbox"/> Electric Funeral	4:53	Black Sabbath	Paranoid	Metal	★★★★★	22
<input checked="" type="checkbox"/> Hand of Doom	7:10	Black Sabbath	Paranoid	Metal	★★★★★	23
<input checked="" type="checkbox"/> Rat Salad	2:30	Black Sabbath	Paranoid	Metal	★★★★★	31
<input checked="" type="checkbox"/> Jack the Stripper/Fairies Wear ...	6:14	Black Sabbath	Paranoid	Metal	★★★★★	24
<input checked="" type="checkbox"/> Bomb Squad (TECH)	3:28	Brent	Brent's Album			1
<input checked="" type="checkbox"/> clay techno	4:36	Brent	Brent's Album			2
<input checked="" type="checkbox"/> Heavy	3:08	Brent	Brent's Album			1
<input checked="" type="checkbox"/> Hi metal man	4:20	Brent	Brent's Album			1
<input checked="" type="checkbox"/> Mistro	2:58	Brent	Brent's Album			1

# For each “piece of info”...

- Is the column an object or an attribute of another object?

Len	Album
Genre	Artist
Rating	Count
Track	

- Avoid duplication

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tin Man	2:20	America	Greatest Hits	Easy Listen...	★★★★★	22

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock		70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...		18
<input checked="" type="checkbox"/> Tin Man	2:20	America	Greatest Hits	Easy Listen...		22

Track  
Album  
Artist  
Genre  
Rating  
Len  
Count

Artist

belongs-to

Album

belongs-to

Track

Rating  
Len  
Count

Genre

belongs-to

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock		70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...		18
<input checked="" type="checkbox"/> Tin Man	2:20	America	Greatest Hits	Easy Listen...		22

# Representing Relationships in a Database

# Database Normalization

- There is \*tons\* of database theory - way too much to understand without excessive predicate calculus
- Do not replicate data - reference data - point at data
- Use integers for keys and for references
- Add a special “key” column to each table which we will make references to. By convention, many programmers call this column “id”

[http://en.wikipedia.org/wiki/Database\\_normalization](http://en.wikipedia.org/wiki/Database_normalization)

# Integer Reference Pattern

We use integers to reference rows in another table

id	name
Filter	Filter
1	Led Zepplin
2	AC/DC

Artist

id	artist_id	title
Filter	Filter	Filter
1	2	Who Made Who
2	1	IV

Album

# Three Kinds of Keys

- Primary key - generally an integer auto-increment field
- Logical key - What the outside world uses for lookup
- Foreign key - generally an integer key pointing to a row in another table



# Key Rules

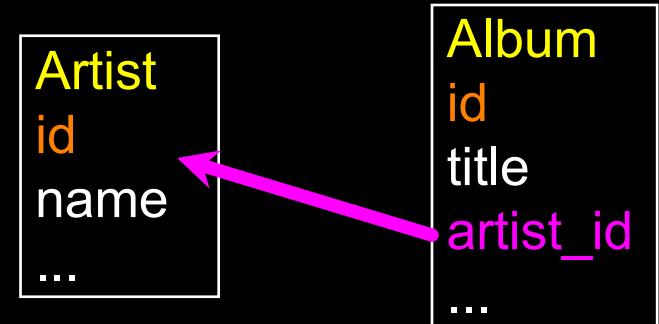
## Best practices

- Never use your **logical key** as the **primary key**
- **Logical keys** can and do change, albeit slowly
- **Relationships** that are based on matching string fields are less efficient than integers

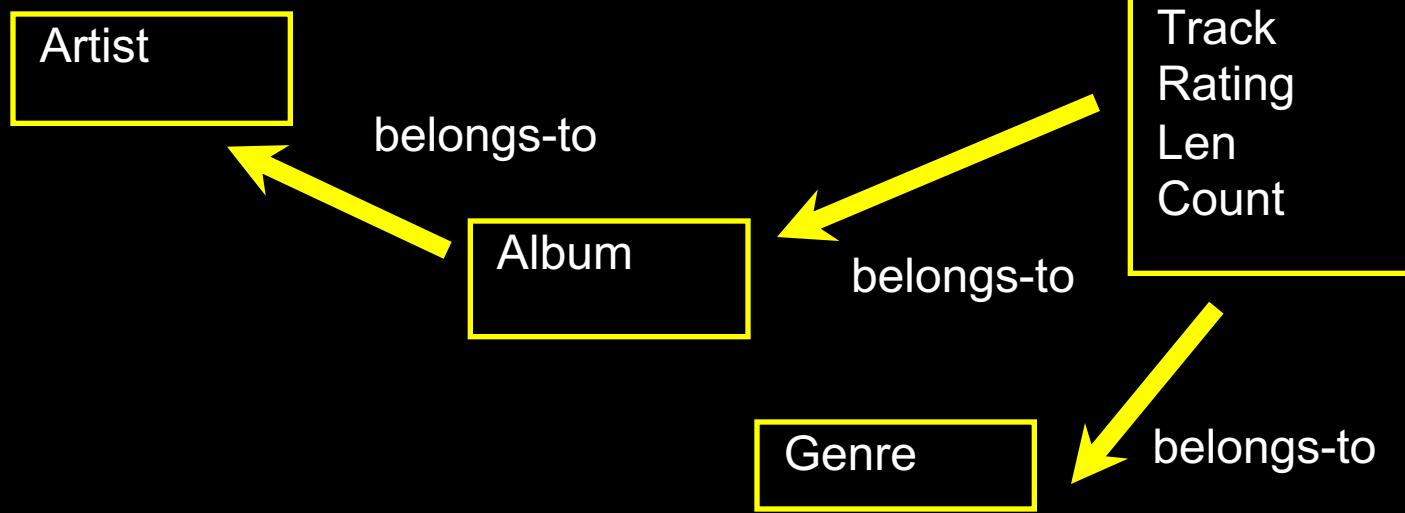
User  
id  
login  
password  
name  
email  
created\_at  
modified\_at  
login\_at

# Foreign Keys

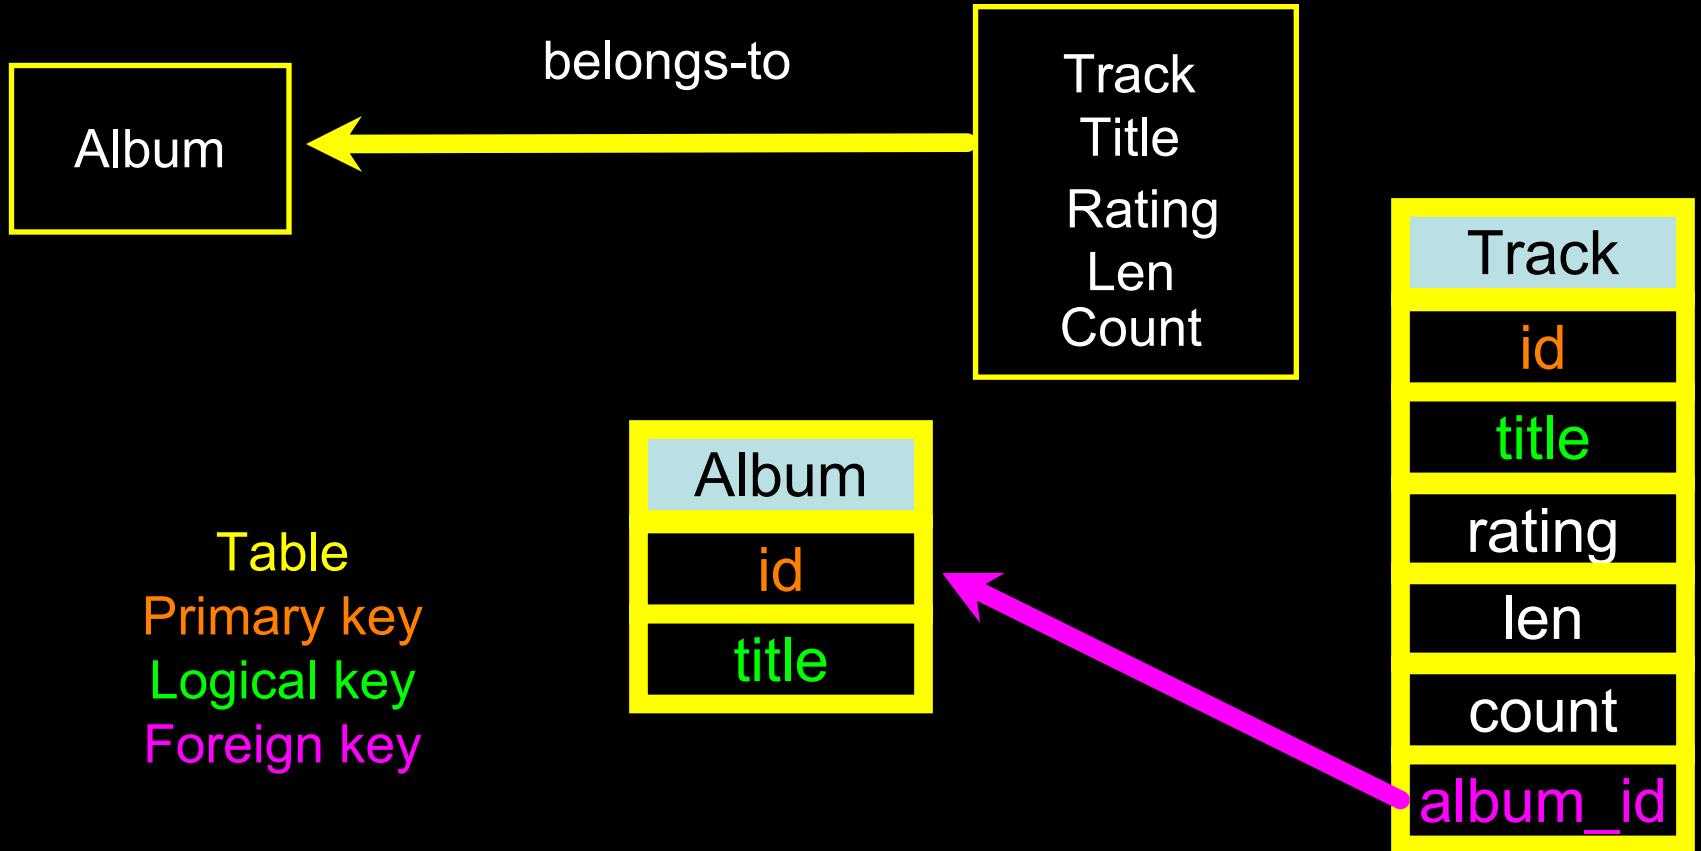
- A **foreign key** is when a table has a column that contains a key which points to the **primary key** of another table.
- When all primary keys are integers, then all foreign keys are integers - this is good - very good
- Good naming convention:  
`<other_table>_id` e.g., `artist_id`



# Relationship Building (in tables)



✓ Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
✓ Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
✓ For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
✓ Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
✓ Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
✓ Tin Man	2:20	America	Greatest Hits	Easy Listen...	★★★★★	22



Artist	
id	
name	

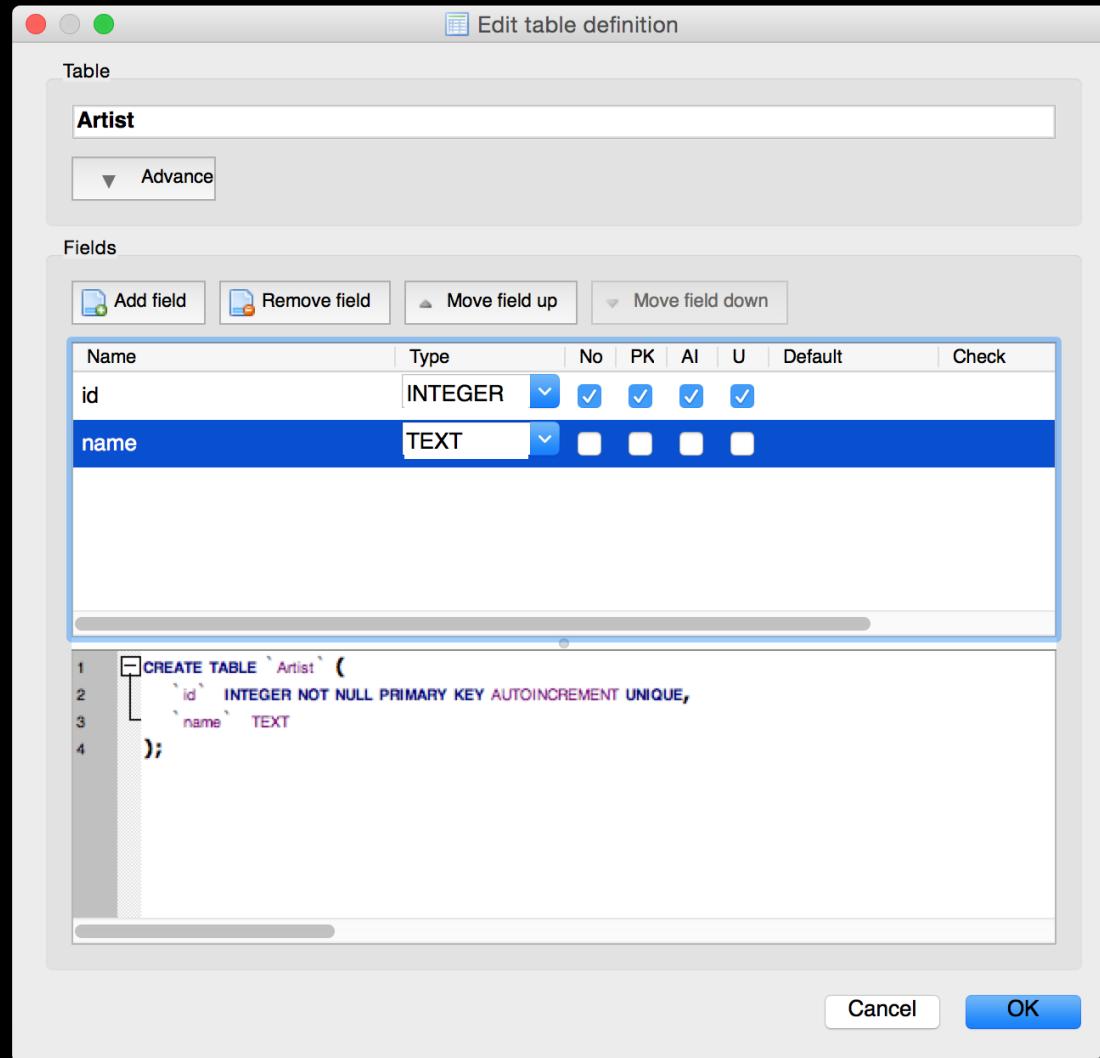
Album	
id	
title	
artist_id	

Track	
id	
title	
rating	
len	
count	
album_id	
genre_id	

Genre	
id	
name	

Table  
Primary key  
Logical key  
Foreign key

Naming the Foreign Key  
artist\_id is a convention



Id is the key:

- No: NOT NULL
- PK: Primary Key
- AI: Auto-incremented
- U: Unique values

DB Browser for SQLite - /Users/csev/Desktop/Music

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

SQL1

```
1 CREATE TABLE Genre (
2     id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
3     name TEXT
4 )
5
```

Query executed successfully: CREATE TABLE Genre ( id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE, name TEXT ) (took 0ms)

DB Schema

Name	Type	Schema
Tables (2)		
Artist	CREATE TABLE `Artist` ( `id` ... )	
id	INTEGER	`id` INTEGER NOT NULL ...
name	TEXT	`name` TEXT
sqlite_sequence	CREATE TABLE sqlite_sequence...	
Indices (1)		
sqlite_autoindex_1		
Views (0)		
Triggers (0)		

Plot DB Schema

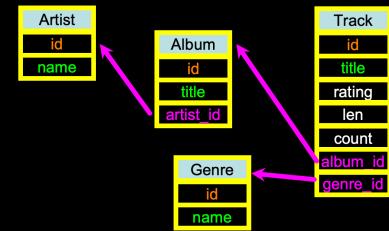
UTF-8

The screenshot shows the DB Browser for SQLite application interface. The main window has tabs for 'Database Structure', 'Browse Data', 'Edit Pragmas', and 'Execute SQL'. The 'Execute SQL' tab is active, displaying an SQL command to create a 'Genre' table with columns 'id' and 'name'. Below the SQL input field, a message indicates the query was executed successfully. To the right, the 'DB Schema' tab is open, showing the created 'Artist' table and its columns ('id' and 'name'), along with other system tables like 'sqlite\_sequence' and indices.

```
CREATE TABLE Genre (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    name    TEXT
)
```

```
CREATE TABLE Album (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    artist_id  INTEGER,
    title    TEXT
)
```

```
CREATE TABLE Track (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    title  TEXT,
    album_id  INTEGER,
    genre_id  INTEGER,
    len     INTEGER, rating INTEGER, count INTEGER
)
```



Database Structure    Browse Data    Edit Pragmas    Execute SQL

Create Table    Modify Table    Delete Table

Name	Type	Schema
<b>Tables (5)</b>		
Album		<pre>CREATE TABLE "Album" (     `id` INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,     `artist_id` INTEGER,     `title` TEXT )</pre>
Artist		<pre>CREATE TABLE 'Artist' (     `id` INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,     `name` TEXT )</pre>
Genre		<pre>CREATE TABLE Genre (     id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,     name TEXT )</pre>
Track		<pre>CREATE TABLE Track (     id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,     title TEXT,     album_id INTEGER,     genre_id INTEGER,     len INTEGER,     rating INTEGER,     count INTEGER )</pre>

Album    Artist    Genre    Track

DB Schema

Name	Schema
<b>Tables (5)</b>	
Album	<pre>CREATE TABLE "Album" (     `id` INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,</pre>
Artist	<pre>CREATE TABLE 'Artist' (     `id` INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,</pre>
Genre	<pre>CREATE TABLE Genre (     id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,</pre>
Track	<pre>CREATE TABLE Track (     id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT,</pre>
sqlite_sequence	<pre>CREATE TABLE sqlite_sequence(name,seq)</pre>
<b>Indices (4)</b>	
sqlite_autoindex_Album_1	
sqlite_autoindex_Album_2	

Plot    DB Schema

UTF-8

DB Browser for SQLite - /Users/csev/Desktop/Music

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

SQL 1

```
1 Insert into Artist (name) values ('AC/DC')
2
```

Query executed successfully: insert into Artist (name) values ('AC/DC') (took 0ms)

DB Schema

Name Schema

Tables (5)

- Album`CREATE TABLE Album (id INTEGER NOT NULL P...)`
- Artist`CREATE TABLE `Artist` (id INTEGER NOT NULL P...`
  - id`'id' INTEGER NOT NULL PRIMARY KEY AUTOINC...`
  - name`'name' TEXT`
- Genre`CREATE TABLE Genre (id INTEGER NOT NULL P...)`
- Track`CREATE TABLE Track (id INTEGER NOT NULL P...)`

sqlite\_sequence`CREATE TABLE sqlite_sequence(name,seq)`

Indices (4)

- sqlite\_auto...
- sqlite\_auto...
- sqlite\_auto...
- sqlite\_auto...

Views (0)

Triggers (0)

Plot DB Schema

UTF-8

The screenshot shows the DB Browser for SQLite interface. The left pane contains a SQL editor with the query `Insert into Artist (name) values ('AC/DC')`. Below it, a message indicates the query was executed successfully. The right pane displays the database schema with five tables: Album, Artist, Genre, Track, and sqlite\_sequence. The Artist table is currently selected. The schema for the Artist table includes columns for id (primary key, auto-increment) and name. There are also four indices for the sqlite\_sequence table and no views or triggers.

insert into Artist (name) values ('Led Zepplin')  
insert into Artist (name) values ('AC/DC')

DB Browser for SQLite - /Users/csev/Desktop/Music

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

SQL 1

```
1 insert into Artist (name) values ('AC/DC')
2
```

Query executed successfully: insert into Artist (name) values ('AC/DC') (took 0ms)

Artist

id	name
1	Led Zepplin
2	AC/DC

Tables (5)

- Album
- Artist
- Genre
- Track
- sqlite\_sequence

Indices (4)

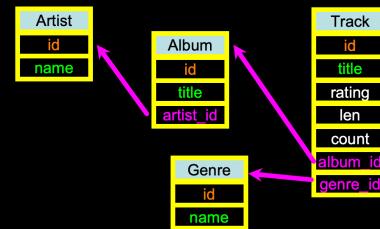
- sqlite\_autoindex\_Album\_1
- sqlite\_autoindex\_Artist\_1
- sqlite\_autoindex\_Genre\_1
- sqlite\_autoindex\_Track\_1

Views (0)

Triggers (0)

DB Schema

insert into Artist (name) values ('Led Zepplin')  
 insert into Artist (name) values ('AC/DC')



DB Browser for SQLite - /Users/csev/Desktop/Music

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

Table: Genre New Record Delete Record

	id	name
1	1	Rock
2	2	Metal

Filter Filter

1 - 2 of 2 Go to: 1

DB Schema

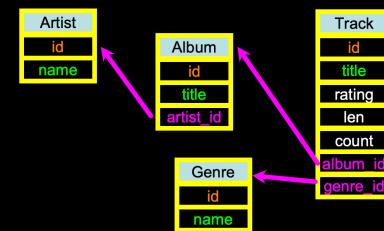
Name Schema

- Tables (5)
  - Album: CREATE TABLE Album (id INTEGER NOT NULL PRIMARY KEY)
  - Artist: CREATE TABLE `Artist` (id INTEGER NOT NULL PRIMARY KEY)
  - Genre: CREATE TABLE Genre (id INTEGER NOT NULL PRIMARY KEY)
  - Track: CREATE TABLE Track (id INTEGER NOT NULL PRIMARY KEY)
  - sqlite\_sequence: CREATE TABLE sqlite\_sequence(name,seq)
- Indices (4)
  - sqlite\_auto...
  - sqlite\_auto...
  - sqlite\_auto...
  - sqlite\_auto...
- Views (0)
- Triggers (0)

Plot DB Schema

UTF-8

insert into Genre (name) values ('Rock')  
 insert into Genre (name) values ('Metal')



DB Browser for SQLite - /Users/csev/Desktop/Music

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

Table: Album New Record Delete Record

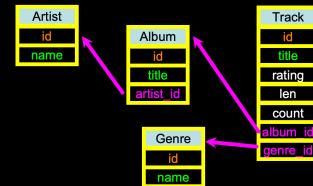
	id	artist_id	title
1	1	2	Who Made Who
2	2	1	IV

< < 1 - 2 of 2 > >| Go to: 1

UTF-8

```
graph LR; Artist[Artist] --> Album[Album]; Artist --> Genre[Genre]; Album --> Track[Track];
```

```
insert into Album (title, artist_id) values ('Who Made Who', 2)
insert into Album (title, artist_id) values ('IV', 1)
```



```

insert into Track (title, rating, len, count, album_id, genre_id)
values ('Black Dog', 5, 297, 0, 2, 1)
insert into Track (title, rating, len, count, album_id, genre_id)
values ('Stairway', 5, 482, 0, 2, 1)
insert into Track (title, rating, len, count, album_id, genre_id)
values ('About to Rock', 5, 313, 0, 1, 2)
insert into Track (title, rating, len, count, album_id, genre_id)
values ('Who Made Who', 5, 207, 0, 1, 2)
  
```

	<b>id</b>	<b>title</b>	<b>album_id</b>	<b>genre_id</b>	<b>len</b>	<b>rating</b>	<b>count</b>
	Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	1	Black Dog	2	1	297	5	0
2	2	Stairway	2	1	482	5	0
3	3	About to Rock	1	2	313	5	0
4	4	Who Made Who	1	2	207	5	0

Track

	id	title	album_id	genre_id	len	rating	count
Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	Black Dog	2	1	297	5	0	
2	Stairway	2	1	482	5	0	
3	About to Rock	1	2	313	5	0	
4	Who Made Who	1	2	207	5	0	

Album

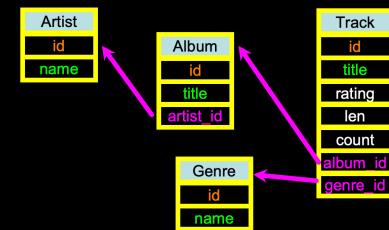
	id	artist_id	title
Filter	Filter	Filter	Filter
1	2		Who Made Who
2	1		IV

Artist

	id	name
Filter	Filter	Filter
1		Led Zeppelin
2		AC/DC

	id	name
Filter	Filter	Filter
1		Rock
2		Metal

Genre



# Declaring a Foreign Key

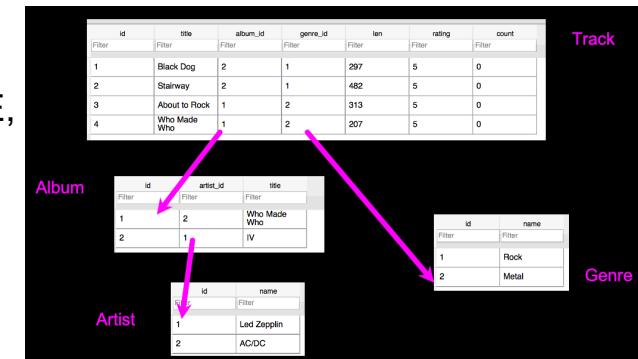
```
CREATE TABLE Artist (
    id INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL UNIQUE,
    name TEXT)
```

```
CREATE TABLE Album (
    id INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL UNIQUE,
    title TEXT,
    artist_id INTEGER,
    FOREIGN KEY(artist_id) REFERENCES Artist(id))
```

```
CREATE TABLE Genre (
    id INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL UNIQUE,
    name TEXT)
```

```
CREATE TABLE Track (
    id INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL UNIQUE,
    title TEXT, rating INTEGER, len INTEGER, count INTEGER,
    album_id INTEGER,
    genre_id INTEGER,
    FOREIGN KEY(album_id) REFERENCES Album(id),
    FOREIGN KEY(genre_id) REFERENCES Genre(id))
```

Database will enforce that *artist\_id* corresponds to a valid *id* in table *Artist*.



# Relate Data Across Tables Using Join

[http://en.wikipedia.org/wiki/Join\\_\(SQL\)](http://en.wikipedia.org/wiki/Join_(SQL))

# Relational Power

- By removing the replicated data and replacing it with references to a single copy of each bit of data we build a “**web**” of information that the relational database can read through very quickly - even for very large amounts of data
- Often when you want some data it comes from a number of tables linked by these **foreign keys**

# The JOIN Operation

- The JOIN operation **links data across several tables** as part of a query
- You must tell the JOIN **how to use the keys** that make the connection between the tables using an **ON clause**

<b>id</b>	<b>artist_id</b>	<b>title</b>
Filter	Filter	Filter
1	2	Who Made Who
2	1	IV

Album

	<b>title</b>	<b>name</b>
1	Who Made Who	AC/DC
2	IV	Led Zepplin

Artist

<b>id</b>	<b>name</b>
Filter	Filter
1	Led Zepplin
2	AC/DC

SELECT Album.title, Artist.name FROM Album JOIN Artist ON Album.artist\_id = Artist.id

What we want  
to see

The tables that  
hold the data

How the tables  
are linked

<b>id</b>	<b>artist_id</b>	<b>title</b>	<b>id</b>	<b>name</b>
Filter	Filter	Filter	Filter	Filter
1	2	Who Made Who	1	Led Zepplin
2	1	IV	2	AC/DC

	<b>title</b>	<b>artist_id</b>	<b>id</b>	<b>name</b>
1	Who Made Who	2	2	AC/DC
2	IV	1	1	Led Zepplin

SELECT Album.title, Album.artist\_id, Artist.id, Artist.name  
 FROM Album JOIN Artist ON Album.artist\_id = Artist.id

	title	genre_id	id	name
1	Black Dog	1	1	Rock
2	Black Dog	1	2	Metal
3	Stairway	1	1	Rock
4	Stairway	1	2	Metal
5	About to Rock	2	1	Rock
6	About to Rock	2	2	Metal
7	Who Made Who	2	1	Rock
8	Who Made Who	2	2	Metal

```
SELECT Track.title,  
       Track.genre_id,  
       Genre.id, Genre.name  
FROM Track JOIN Genre
```

Joining two tables without an  
**ON** clause gives all possible  
combinations of rows  
(aka the cartesian product)

	title	name
id		
1	Black Dog	Rock
2	Stairway	Rock
3	About to Rock	Metal
4	Who Made Who	Metal

id	title	album_id	genre_id	len	rating	count
Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	Black Dog	2	1	297	5	0
2	Stairway	2	1	482	5	0
3	About to Rock	1	2	313	5	0
4	Who Made Who	1	2	207	5	0

id	name
Filter	Filter
1	Rock
2	Metal

select Track.title, Genre.name from Track join Genre on Track.genre\_id = Genre.id

What we want  
to see

The tables that  
hold the data

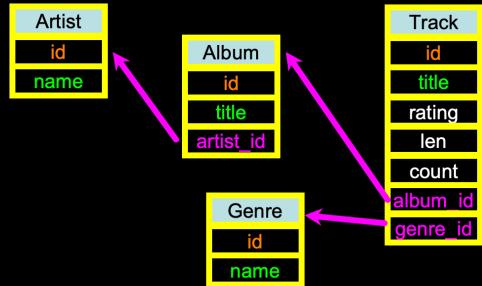
How the tables  
are linked

```

SELECT Track.title, Artist.name, Album.title, Genre.name
FROM Track JOIN Genre JOIN Album JOIN Artist
ON Track.genre_id = Genre.id and
Track.album_id = Album.id and
Album.artist_id = Artist.id

```

	title	name	title	name
1	Black Dog	Led Zepplin	IV	Rock
2	Stairway	Led Zepplin	IV	Rock
3	About to Rock	AC/DC	Who Made Who	Metal
4	Who Made Who	AC/DC	Who Made Who	Metal



What we want to see

The tables which hold  
the data

How the tables are  
linked

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61		
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70		
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56		
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61		
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31		
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23		
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18		
<input checked="" type="checkbox"/> Tin Man	3:30	America	Greatest Hits	Easy Listen...	★★★★★	23		
<input checked="" type="checkbox"/> Sister Golden Hair	3:22	America	Greatest Hits	Easy Listen...	★★★★★	24		
<input checked="" type="checkbox"/> Track 01	4:22	Billy Price	Danger Zone	Blues/R&B	★★★★★	26		
<input checked="" type="checkbox"/> Track 02	2:45	Billy Price	Danger Zone	Blues/R&B	★★★★★	18		
<input checked="" type="checkbox"/> Track 03	3:26	Billy Price	Danger Zone	Blues/R&B	★★★★★	18		
<input checked="" type="checkbox"/> Track 04	4:17	Billy Price	Danger Zone	Blues/R&B	★★★★★	18		
<input checked="" type="checkbox"/> Track 05	3:50	Billy Price	Danger Zone	Blues/R&B	★★★★★	18		
<input checked="" type="checkbox"/> War Pigs/Luke's Wall	7:58	Black Sabbath	Pa	1	Black Dog	Led Zepplin	IV	Rock
<input checked="" type="checkbox"/> Paranoid	2:53	Black Sabbath	Pa	2	Stairway	Led Zepplin	IV	Rock
<input checked="" type="checkbox"/> Planet Caravan	4:35	Black Sabbath	Pa	3	About to Rock	AC/DC	Who Made Who	Metal
<input checked="" type="checkbox"/> Iron Man	5:59	Black Sabbath	Pa	4	Who Made Who	AC/DC	Who Made Who	Metal
<input checked="" type="checkbox"/> Electric Funeral	4:53	Black Sabbath	Pa					
<input checked="" type="checkbox"/> Hand of Doom	7:10	Black Sabbath	Pa					
<input checked="" type="checkbox"/> Rat Salad	2:30	Black Sabbath	Pa					
<input checked="" type="checkbox"/> Jack the Stripper/Fairies Wear ...	6:14	Black Sabbath	Pa					
<input checked="" type="checkbox"/> Bomb Squad (TECH)	3:28	Brent	Br					
<input checked="" type="checkbox"/> clay techno	4:36	Brent	Br					
<input checked="" type="checkbox"/> Heavy	3:08	Brent	Br					
<input checked="" type="checkbox"/> Hi metal man	4:20	Brent	Brent's Album				1	
<input checked="" type="checkbox"/> Mistro	2:58	Brent	Brent's Album				1	

		title	name	title	name
1		Black Dog	Led Zepplin	IV	Rock
2		Stairway	Led Zepplin	IV	Rock
3		About to Rock	AC/DC	Who Made Who	Metal
4		Who Made Who	AC/DC	Who Made Who	Metal

# Alternative Notations for Joins

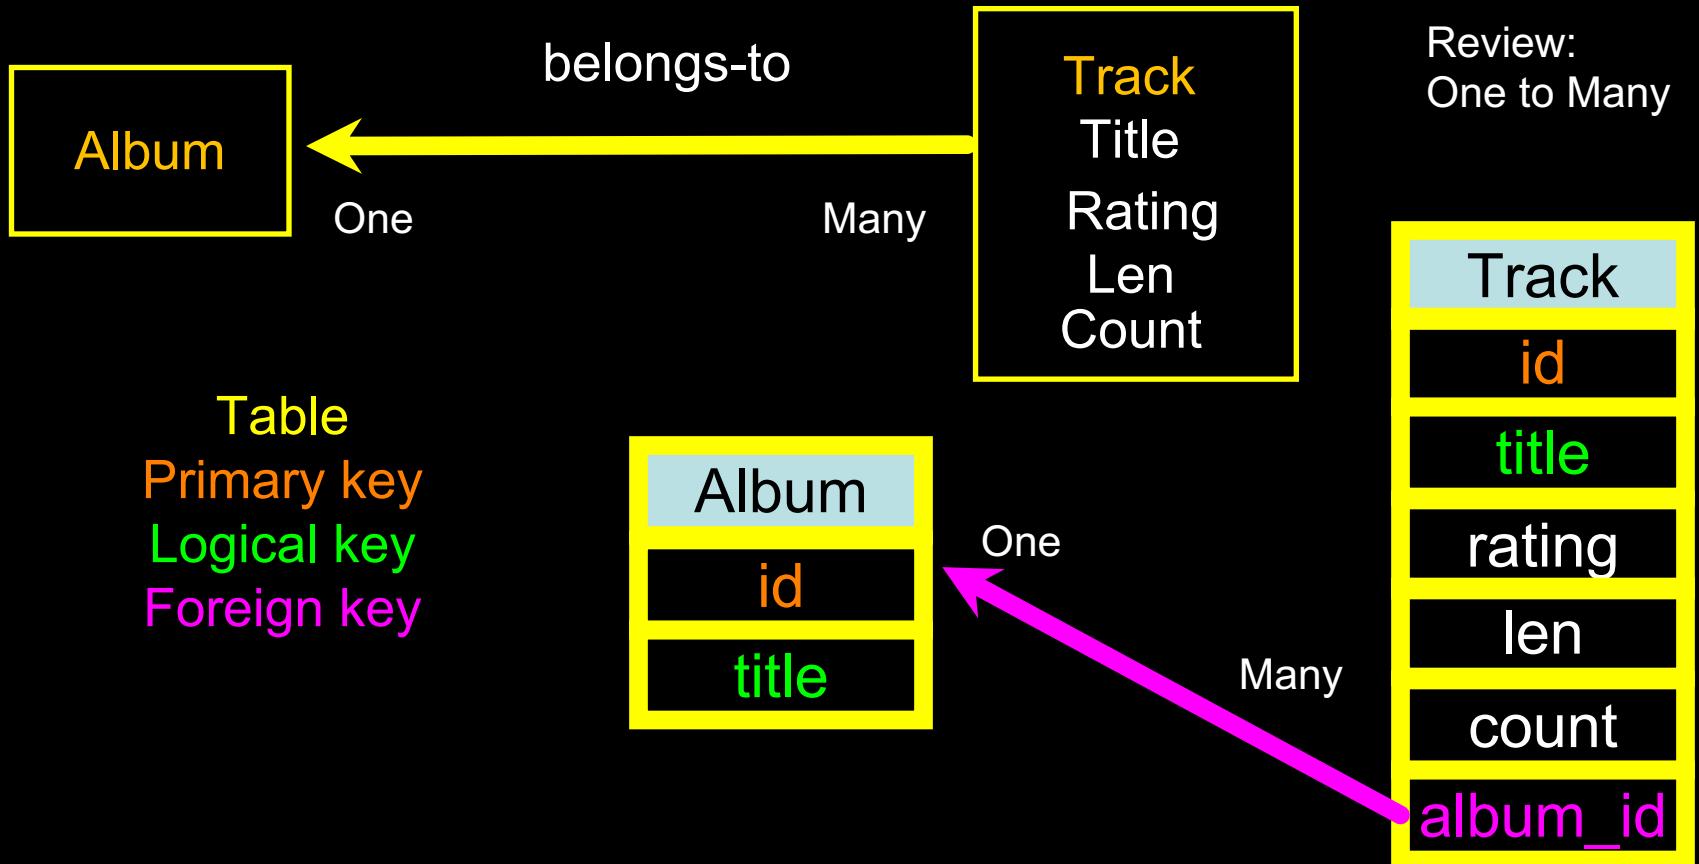
```
SELECT Track.title, Artist.name, Album.title, Genre.name  
FROM Track JOIN Genre JOIN Album JOIN Artist  
ON Track.genre_id = Genre.id and  
Track.album_id = Album.id and  
Album.artist_id = Artist.id
```

What we want to see  
The tables which hold the data  
How the tables are linked

```
SELECT Track.title, Artist.name, Album.title, Genre.name  
FROM Track, Genre, Album, Artist  
WHERE Track.genre_id = Genre.id and  
Track.album_id = Album.id and  
Album.artist_id = Artist.id
```

# Many-To-Many Relationships

[https://en.wikipedia.org/wiki/Many-to-many\\_\(data\\_model\)](https://en.wikipedia.org/wiki/Many-to-many_(data_model))



id		name	
Filter	Filter	Filter	Filter
1		Rock	
2		Metal	

One



One

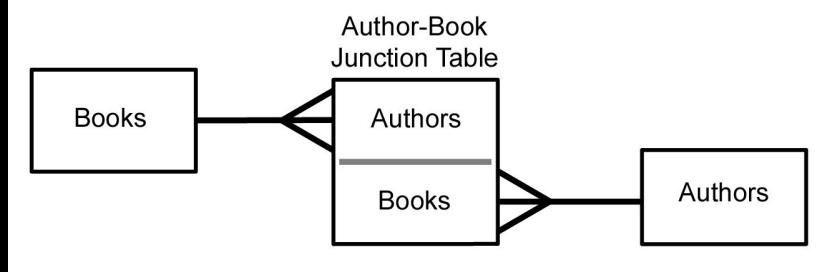
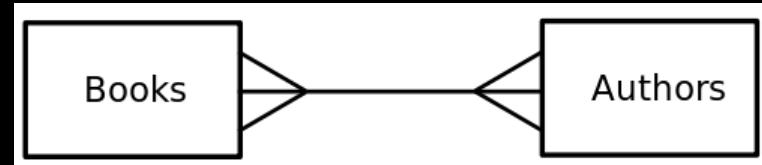
Many

Many

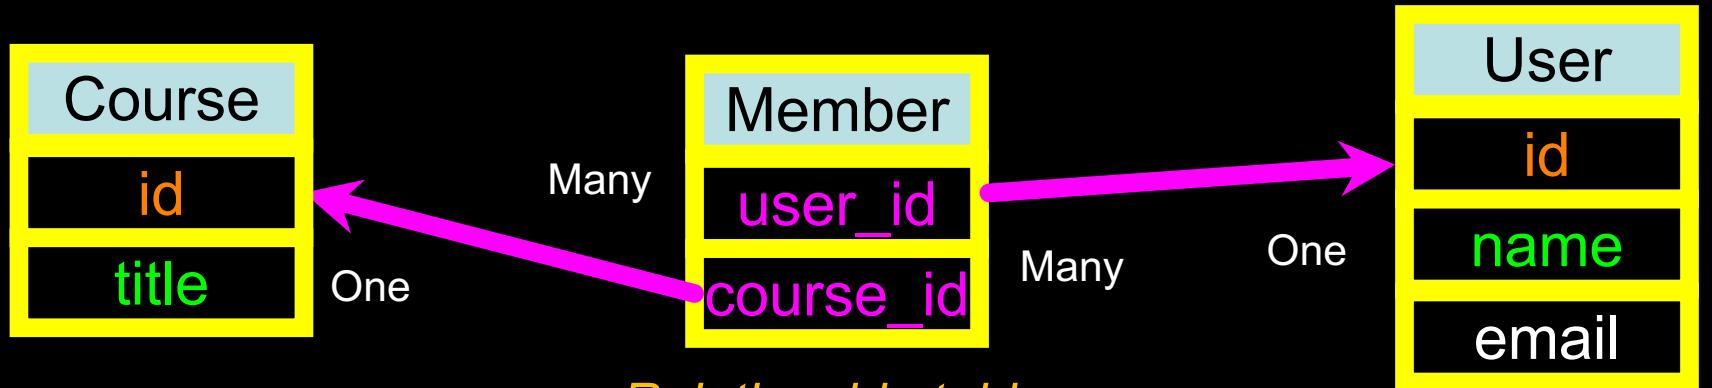
id		title		album_id	genre_id	len	rating	count
Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter
1		Black Dog		2	1	297	5	0
2		Stairway		2	1	482	5	0
3		About to Rock		1	2	313	5	0
4		Who Made Who		1	2	207	5	0

# Many to Many

- Sometimes we need to model a relationship that is many-to-many
- We need to add a "connection" table with two foreign keys
- There is usually no separate primary key



[https://en.wikipedia.org/wiki/Many-to-many\\_\(data\\_model\)](https://en.wikipedia.org/wiki/Many-to-many_(data_model))



*Entity table*

*Relationship table  
to relate courses and users  
(with two foreign keys)*

*Entity table*

```
CREATE TABLE User (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    name    TEXT UNIQUE,
    email   TEXT
)
```

```
CREATE TABLE Course (
    id      INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE,
    title   TEXT UNIQUE
)
```

```
CREATE TABLE Member (
    user_id      INTEGER,
    course_id    INTEGER,
    role         INTEGER,
    PRIMARY KEY (user_id, course_id)
)
```

*In addition to the foreign keys of the entities that the connection (relationship) table links, it can have additional attributes, like role*

DB Browser for SQLite - /Users/csev/Desktop/si502\_database

New Database Open Database Write Changes Revert Changes

Database Structure Browse Data Edit Pragmas Execute SQL

Create Table Modify Table Delete Table

Name	Type	Schema
Tables (4)		
Course		CREATE TABLE Course ( id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE, title TEXT )
Member		CREATE TABLE Member ( user_id INTEGER, course_id INTEGER, PRIMARY KEY (user_id, course_id) )
User		CREATE TABLE User ( id INTEGER NOT NULL PRIMARY KEY AUTOINCREMENT UNIQUE, name TEXT, email TEXT )
sqlite_sequence		CREATE TABLE sqlite_sequence(name,seq)
Indices (3)		
sqlite_autoindex_Course_1		
sqlite_autoindex_Member_1		
sqlite_autoindex_User_1		
Views (0)		
Triggers (0)		

UTF-8

# Insert Users and Courses

```
INSERT INTO User (name, email) VALUES ('Jane', 'jane@tsugi.org');  
INSERT INTO User (name, email) VALUES ('Ed', 'ed@tsugi.org');  
INSERT INTO User (name, email) VALUES ('Sue', 'sue@tsugi.org');
```

```
INSERT INTO Course (title) VALUES ('Python');  
INSERT INTO Course (title) VALUES ('SQL');  
INSERT INTO Course (title) VALUES ('PHP');
```

The image shows two separate windows of the DB Browser for SQLite application. Both windows have a title bar 'DB Browser for SQLite - /Users/csev/Desktop/si502\_database' and a menu bar with 'File', 'Edit', 'View', 'Tools', 'Help'. The left window displays the 'User' table with three rows:

	id	name	email
1	1	Jane	jane@tsugi.org
2	2	Ed	ed@tsugi.org
3	3	Sue	sue@tsugi.org

The right window displays the 'Course' table with three rows:

	id	title
1	1	Python
2	2	SQL
3	3	PHP

## User

id	name	email
Filter	Filter	Filter
1	Jane	jane@tsugi.org
2	Ed	ed@tsugi.org
3	Sue	sue@tsugi.org

## Course

id	title
Filter	Filter
1	Python
2	SQL
3	PHP

Table: Member

user_id	course_id	role
Filter	Filter	Filter
1	1	1
2	2	0
3	3	0
4	1	0
5	2	1
6	2	1
7	3	0

```
INSERT INTO Member (user_id, course_id, role) VALUES (1, 1, 1);
INSERT INTO Member (user_id, course_id, role) VALUES (2, 1, 0);
INSERT INTO Member (user_id, course_id, role) VALUES (3, 1, 0);

INSERT INTO Member (user_id, course_id, role) VALUES (1, 2, 0);
INSERT INTO Member (user_id, course_id, role) VALUES (2, 2, 1);

INSERT INTO Member (user_id, course_id, role) VALUES (2, 3, 1);
INSERT INTO Member (user_id, course_id, role) VALUES (3, 3, 0);
```

id	name	email
Filter	Filter	Filter
1	Jane	jane@tsugi.org
2	Ed	ed@tsugi.org
3	Sue	sue@tsugi.org

user_id	course_id	role
Filter	Filter	Filter
1	1	1
2	1	0
3	1	0
1	2	0
2	2	1
2	3	1
3	3	0

id	title
Filter	Filter
1	Python
2	SQL
3	PHP

```
SELECT User.name, Member.role, Course.title
FROM User JOIN Member JOIN Course
ON Member.user_id = User.id AND
Member.course_id = Course.id
ORDER BY Course.title, Member.role DESC, User.name
```

	name	role	title
2	Sue	0	PHP
3	Jane	1	Python
4	Ed	0	Python
5	Sue	0	Python
6	Ed	1	SQL

# Run Code

Jupyter Notebook:  
week11-multi-tables-db.ipynb

Go to SQLite:  
Run Agents-customers-orders-queries.sql

# SQL Data Definition Commands

COMMAND OR OPTION	DESCRIPTION
CREATE SCHEMA AUTHORIZATION	Creates a database schema
CREATE TABLE	Creates a new table in the user's database schema
NOT NULL	Ensures that a column will not have null values
UNIQUE	Ensures that a column will not have duplicate values
PRIMARY KEY	Defines a primary key for a table
FOREIGN KEY	Defines a foreign key for a table
DEFAULT	Defines a default value for a column (when no value is given)
CHECK	Validates data in an attribute
CREATE INDEX	Creates an index for a table
CREATE VIEW	Creates a dynamic subset of rows and columns from one or more tables (see Chapter 8, Advanced SQL)
ALTER TABLE	Modifies a table's definition (adds, modifies, or deletes attributes or constraints)
CREATE TABLE AS	Creates a new table based on a query in the user's database schema
DROP TABLE	Permanently deletes a table (and its data)
DROP INDEX	Permanently deletes an index
DROP VIEW	Permanently deletes a view

# SQL Data Manipulation Commands

COMMAND OR OPTION	DESCRIPTION
INSERT	Inserts row(s) into a table
SELECT	Selects attributes from rows in one or more tables or views
WHERE	Restricts the selection of rows based on a conditional expression
GROUP BY	Groups the selected rows based on one or more attributes
HAVING	Restricts the selection of grouped rows based on a condition
ORDER BY	Orders the selected rows based on one or more attributes
UPDATE	Modifies an attribute's values in one or more table's rows
DELETE	Deletes one or more rows from a table
COMMIT	Permanently saves data changes
ROLLBACK	Restores data to their original values
<b>Comparison operators</b>	
=, <, >, <=, >=, <>	Used in conditional expressions
<b>Logical operators</b>	
AND/OR/NOT	Used in conditional expressions
<b>Special operators</b>	
BETWEEN	Checks whether an attribute value is within a range
IS NULL	Checks whether an attribute value is null
LIKE	Checks whether an attribute value matches a given string pattern
IN	Checks whether an attribute value matches any value within a value list
EXISTS	Checks whether a subquery returns any rows
DISTINCT	Limits values to unique values
<b>Aggregate functions</b>	
COUNT	Returns the number of rows with non-null values for a given column
MIN	Returns the minimum attribute value found in a given column
MAX	Returns the maximum attribute value found in a given column
SUM	Returns the sum of all values for a given column
AVG	Returns the average of all values for a given column

# Relational Modeling Enables Speed

- Complexity makes speed possible and allows you to get very fast results as the data size grows
- By normalizing the data and linking it with integer keys, the overall amount of data which the relational database must scan is far lower than if the data were simply flattened out
- It might seem like a tradeoff - spend some time designing your database so it continues to be fast when your application is a success

# Additional SQL Topics

- **Indexes (Indexes)** improve data access/query performance
- **Constraints** on data - (cannot be NULL, etc..)
- **Transactions** - allow SQL operations to be grouped and done as a unit

# Indices (Indexes)

## Purpose: Speed up data retrieval

Require some time for indexing, more space, but give great query speedups  
(I once created a table with 3 billion rows in Postgres, worked great after indexing!)

```
CREATE TABLE contacts (
    first_name text NOT NULL,
    last_name text NOT NULL,
    email text NOT NULL
);
```

```
CREATE INDEX idx_contacts_email
ON contacts (email)
```

### Faster due to index:

```
SELECT * FROM contacts WHERE email = 'ulf@isi.edu'
```

# Indexes (multi-column)

```
CREATE TABLE contacts (
    first_name text NOT NULL,
    last_name text NOT NULL,
    email text NOT NULL
);
```

```
CREATE UNIQUE INDEX idx_contacts_name
    ON contacts (first_name, last_name)
```

**Faster due to index:**

```
SELECT * FROM contacts
WHERE first_name = 'Ulf' AND last_name = 'Hermjakob'
```

CREATE UNIQUE INDEX means that duplicate index entries are not allowed. Any attempt to insert a duplicate entry will result in an error.

# Transactions

**Purpose: Transactions ensure data integrity**

An important strength of a database system

**Transactions of SQLite (and other DBMSs) are ACID compliant**

- **Atomic:** Transaction cannot be broken down into smaller parts.  
All operations in a transaction must be completed (else abort).
- **Consistent:** Change DB from one valid state to another
- **Isolation:** Pending transaction must be invisible to other transactions.  
Data used during transaction cannot be used by another transaction
- **Durable:** A committed transaction must be permanent
  - even in the face of power failures or system crashes

# Transaction Example: Bank Transfer

```
BEGIN TRANSACTION;  
  
UPDATE accounts  
    SET balance = balance - 1000  
    WHERE account_no = 100;  
  
UPDATE accounts  
    SET balance = balance + 1000  
    WHERE account_no = 200;  
  
INSERT INTO account_changes(account_no,flag,amount,changed_at)  
VALUES(100,'-',1000,datetime('now'));  
  
INSERT INTO account_changes(account_no,flag,amount,changed_at)  
VALUES(200,'+',1000,datetime('now'));  
  
COMMIT
```

Use **ROLLBACK** command  
to rollback uncommitted  
transactions.

# Summary

- Relational databases allow us to **scale** to very large amounts of data
- The key is to have **one copy of any data element** and use relations and joins to link the data to multiple places
- This greatly **reduces the amount of data which much be scanned** when doing complex operations across large amounts of data
- Database and SQL design is a bit of an **art form**